

Lecture 03.4 Pandas

August 31, 2023

1 Intro to Pandas

Duncan Callaway

My objective in this notebook is to review enough of the basics to give people a clear sense of * what a pandas **data frame** is * what it can contain * how to access the information in it

1.1 Pandas references

Introductory:

- [Getting started with Python for research](#), a gentle introduction to Python in data-intensive research.
- [A Whirlwind Tour of Python](#), by Jake VanderPlas, another quick Python intro (with notebooks).

Core Pandas/Data Science books:

- [The Python Data Science Handbook](#), by Jake VanderPlas.
- [Python for Data Analysis, 2nd Edition](#), by Wes McKinney, creator of Pandas. [Companion Notebooks](#)
- [Effective Pandas](#), a book by Tom Augspurger, core Pandas developer.

Complementary resources:

- [An introduction to “Data Science”](#), a collection of Notebooks by BIDS’ Stéfan Van der Walt.
- [Effective Computation in Physics](#), by Kathryn D. Huff; Anthony Scopatz. [Notebooks to accompany the book](#). Don’t be fooled by the title, it’s a great book on modern computational practices with very little that’s physics-specific.

1.2 Introduction

In Data8 you used the `tables` library to organize and manipulate data. I’ve not used tables but it has been described to me as a ‘light’ version of pandas. So this will be somewhat familiar to you.

Pandas has several features that make it immediately better than numpy for organizing data. 1. You can have different data types (string, int, float) in each column 2. You can label columns (headers) 3. You can label rows (index)

We call the data structure that holds all these things together a **data frame**

Ambiguity alert: up to now we’ve talked about indexing to access individual elements of a data structure. The index in pandas has a slightly more specific meaning, in that it references the rows of the data frame. Pandas documentation talks about “location” or “position” in place using “index” as we did with the other data structures. I’ll try my best to disambiguate, but also rely on context to clarify the meaning of the terms.

At its core, the data frame structure is what draws us to use pandas. But it also has a bunch of great built-in functions you can use to manipulate data once it is loaded in.

```
[1]: import pandas as pd
```

I’m going to define a simple data frame in a way that you can see the connections to existing Python data types and structures.

First let’s define a dict of fruit information:

```
[2]: fruit_info_dict = {'fruit': ['apple', 'banana', 'orange', 'raspberry'],
                        'color': ['red', 'yellow', 'orange', 'pink'],
                        'weight': [120, 150, 250, 15]}
fruit_info_dict
```

```
[2]: {'fruit': ['apple', 'banana', 'orange', 'raspberry'],
      'color': ['red', 'yellow', 'orange', 'pink'],
      'weight': [120, 150, 250, 15]}
```

1.2.1 Q: What do you get if you call the dict of lists with a key?

```
[3]: fruit_info_dict['color']
```

```
[3]: ['red', 'yellow', 'orange', 'pink']
```

Ans: A the list associated with the key

1.2.2 Q: Figure out how to get the weight of a raspberry out of the dict.

```
[4]: fruit_info_dict['weight'][3]
```

```
[4]: 15
```

1.3 Now let’s make a dataframe

It’s just a fancy version of a dict of lists:

```
[5]: fruit_info_df = pd.DataFrame(
    data={'fruit': ['apple', 'banana', 'orange', 'raspberry'],
          'color': ['red', 'yellow', 'orange', 'pink'],
          'weight': [120, 150, 250, 15]}
    )
fruit_info_df
```

```
[5]:      fruit  color  weight
0     apple   red    120
1    banana yellow    150
2     orange orange    250
3  raspberry  pink     15
```

Some notes: 1. You can see that we put the data inside curly brackets much like we do in a dict. 1. We defined it as if fruit, color and weight are keys. 1. In fact, we'll be calling these column names. 2. The pandas data structure is called the data frame.

```
[6]: type(fruit_info_df)
```

```
[6]: pandas.core.frame.DataFrame
```

We can call the values associated with each column name (like the dict key) in much the same way that we did with the dict:

```
[7]: fruit_info_dict['color']
```

```
[7]: ['red', 'yellow', 'orange', 'pink']
```

```
[8]: fruit_info_df['color']
```

```
[8]: 0      red
1    yellow
2    orange
3     pink
Name: color, dtype: object
```

But notice above in the dataframe we don't just have a list of colors. Instead we have something called a **series**. This is a pandas object that is analogous to a numpy series.

```
[9]: type(fruit_info_df['color'])
```

```
[9]: pandas.core.series.Series
```

The series differs from the list in at least one important way: It has numbers directly associated with it that we call the index. (The left column of the Series.)

Note that in the above, the index is just numeric. But as we'll see, we can make it whatever we want. This makes the data frame much more flexible than a list – we can call elements from a sensical index, rather than just a number.

We can also call columns from the data frame as follows:

```
[10]: fruit_info_df.color
```

```
[10]: 0      red
1    yellow
2    orange
3     pink
```

Name: color, dtype: object

But as we'll see soon, there are alternative ways to get access to the elements of the data frame (`.loc` and `.iloc`) that enable us to work with the frame more as we would a numpy array.

First, let me show you how to get into the dataframe to index individual elements if we *don't* use `.loc` or `.iloc`.

In this case you index into the df in a way that looks a bit like indexing into a *list*. That is, you use two sequences of square brackets, the first carrying information about the column and the second information about the row.

```
[11]: fruit_info_df['fruit'][1]
```

```
[11]: 'banana'
```

```
[12]: fruit_info_df.fruit[1]
```

```
[12]: 'banana'
```

Slicing is limited in this case: we can only slice down rows, not columns of the df:

```
[13]: fruit_info_df['fruit'][0:2]
```

```
[13]: 0    apple
      1    banana
      Name: fruit, dtype: object
```

1.3.1 Anatomy of the data frame.

Let's talk a little about the anatomy of the data frame.

We have the following important attributes: 1. Rows 2. Columns 3. Index 4. Column names

The "index" can be numeric, but as we'll see we can also make the indices strings.

1.3.2 Pandas and the CAISO data.

Let's see what's in the current directory. In macos, I type:

```
[14]: !ls
```

```
CAISO_2017to2018.csv      Lecture 03.3 Numpy.pdf
Lecture 03 InClassNotebook.ipynb  Lecture 03.4 Pandas.ipynb
Lecture 03.1 DataTypes Pandas.pptx  Lecture 03.4 Pandas.pdf
Lecture 03.2 DataTypes.ipynb        dataframe_anatomy.png
Lecture 03.2 DataTypes.pdf          from drive
Lecture 03.3 Numpy.ipynb            ~$Lecture 03.1 DataTypes Pandas.pptx
```

But if someone picked up my ipynb and ran it on windows, they'd get an error, since windows doesn't recognize `ls` but rather `dir`.

Instead we can use the `os` library:

```
[15]: import os
      os.listdir()
```

```
[15]: ['Lecture 03.3 Numpy.pdf',
      '.DS_Store',
      '~$Lecture 03.1 DataTypes Pandas.pptx',
      'dataframe_anatomy.png',
      'Lecture 03.2 DataTypes.pdf',
      'Lecture 03.4 Pandas.ipynb',
      'Lecture 03 InClassNotebook.ipynb',
      'Lecture 03.1 DataTypes Pandas.pptx',
      'Lecture 03.3 Numpy.ipynb',
      '.ipynb_checkpoints',
      'from drive',
      'Lecture 03.4 Pandas.pdf',
      'Lecture 03.2 DataTypes.ipynb',
      'CAISO_2017to2018.csv']
```

I've already downloaded a file, "CAISO_2017to2018.csv", that has one year of renewables production data from CAISO.

Let's load that in as a dataframe and take a look at it.

The simplest command is `pd.read_csv`

```
[16]: caiso_data = pd.read_csv('CAISO_2017to2018.csv')
```

Now we can look at the top of the dataframe using the `.head` method.

Note that you need to put parentheses on the end of the call, otherwise python returns the head "object" in a rather ugly form.

```
[17]: caiso_data.head()
```

```
[17]:
```

	Unnamed: 0	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO	WIND	TOTAL	\
0	2017-08-29 00:00:00	1181	340	156	324		1551	
1	2017-08-29 01:00:00	1182	338	156	326		1556	
2	2017-08-29 02:00:00	1183	337	156	337		1325	
3	2017-08-29 03:00:00	1185	339	156	313		1158	
4	2017-08-29 04:00:00	1190	344	156	320		1209	

	SOLAR PV	SOLAR THERMAL
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

Pandas loaded the date and time in as a column and put its own row numbers on the data frame.

As an alternative we can actually make the row labels *equal* to whatever column of data we'd like.

We'll come back to the notion of the index a little later, but for now let's just reload with the date-time as the index:

```
[18]: caiso_data = pd.read_csv('CAISO_2017to2018.csv', index_col=0)
      caiso_data.head()
```

```
[18]:
```

	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO	WIND TOTAL \
2017-08-29 00:00:00	1181	340	156	324	1551
2017-08-29 01:00:00	1182	338	156	326	1556
2017-08-29 02:00:00	1183	337	156	337	1325
2017-08-29 03:00:00	1185	339	156	313	1158
2017-08-29 04:00:00	1190	344	156	320	1209

	SOLAR PV	SOLAR THERMAL
2017-08-29 00:00:00	0	0
2017-08-29 01:00:00	0	0
2017-08-29 02:00:00	0	0
2017-08-29 03:00:00	0	0
2017-08-29 04:00:00	0	0

We can also rename column names to make things prettier

```
[19]: caiso_data = pd.read_csv('CAISO_2017to2018.csv')
      caiso_data.columns
```

```
[19]: Index(['Unnamed: 0', 'GEO THERMAL', 'BIOMASS', 'BIOGAS', 'SMALL HYDRO',
            'WIND TOTAL', 'SOLAR PV', 'SOLAR THERMAL'],
            dtype='object')
```

Note that we can't reassign easily because column and index names lists are immutable. Here is the workaround:

```
[20]: cols = caiso_data.columns.tolist()
      cols[0] = 'Date and time'
      caiso_data.columns = cols
      caiso_data
```

```
[20]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO \
0	2017-08-29 00:00:00	1181	340	156	324
1	2017-08-29 01:00:00	1182	338	156	326
2	2017-08-29 02:00:00	1183	337	156	337
3	2017-08-29 03:00:00	1185	339	156	313
4	2017-08-29 04:00:00	1190	344	156	320
...
8755	2018-08-28 19:00:00	962	332	236	581
8756	2018-08-28 20:00:00	967	336	234	547
8757	2018-08-28 21:00:00	972	336	233	502
8758	2018-08-28 22:00:00	975	333	234	361
8759	2018-08-28 23:00:00	977	333	235	262

	WIND TOTAL	SOLAR PV	SOLAR THERMAL
0	1551	0	0
1	1556	0	0
2	1325	0	0
3	1158	0	0
4	1209	0	0
...
8755	3300	70	24
8756	3468	0	17
8757	3310	0	17
8758	3068	0	0
8759	2921	0	0

[8760 rows x 8 columns]

Ok, that looks a little better for now.

As you can see, all the data are the same type of numeric value – MWh.

In these cases, sometimes it’s natural to “stack” the data.

We could do the stacking with a pandas command, `.stack`

Let’s work with a few things to explore the data frame.

We’ve already learned about `.head` but we can use it slightly differently:

```
[21]: caiso_data.head(2) # the number in the parens tells pandas how many rows
```

```
[21]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO	WIND TOTAL \
0	2017-08-29 00:00:00	1181	340	156	324	1551
1	2017-08-29 01:00:00	1182	338	156	326	1556

	SOLAR PV	SOLAR THERMAL
0	0	0
1	0	0

We can also look at the tail!

```
[22]: caiso_data.tail()
```

```
[22]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO \
8755	2018-08-28 19:00:00	962	332	236	581
8756	2018-08-28 20:00:00	967	336	234	547
8757	2018-08-28 21:00:00	972	336	233	502
8758	2018-08-28 22:00:00	975	333	234	361
8759	2018-08-28 23:00:00	977	333	235	262

	WIND TOTAL	SOLAR PV	SOLAR THERMAL
8755	3300	70	24

8756	3468	0	17
8757	3310	0	17
8758	3068	0	0
8759	2921	0	0

What's the shape of the frame?

```
[23]: caiso_data.shape
```

```
[23]: (8760, 8)
```

The `.shape` method returns a tuple – number of rows and number of columns.

What about the total number of entries? Use the `.size` method:

```
[24]: caiso_data.size
```

```
[24]: 70080
```

We can also summarize all the numeric data:

```
[25]: caiso_data.describe()
```

```
[25]:
```

	GEOTHERMAL	BIOMASS	BIOGAS	SMALL HYDRO	WIND TOTAL \
count	8760.000000	8760.000000	8760.000000	8760.000000	8760.000000
mean	949.228881	329.203311	224.342808	394.902626	1806.004338
std	108.304664	43.243815	24.228372	112.120046	1284.668963
min	468.000000	164.000000	133.000000	148.000000	0.000000
25%	923.000000	301.000000	219.000000	308.000000	598.000000
50%	956.000000	331.000000	234.000000	379.000000	1615.000000
75%	986.000000	363.000000	240.000000	476.250000	2900.250000
max	1230.000000	482.000000	253.000000	681.000000	5006.000000

	SOLAR PV	SOLAR THERMAL
count	8760.000000	8760.000000
mean	2988.774658	133.840753
std	3628.189420	200.374857
min	0.000000	0.000000
25%	0.000000	0.000000
50%	175.500000	0.000000
75%	6700.250000	261.000000
max	10102.000000	679.000000

1.3.3 Q: In class. Extract the first 5 entries of BIOGAS and SOLAR PV

```
[26]: caiso_data[['BIOGAS', 'SOLAR PV']][0:4]
```

```
[26]:
```

	BIOGAS	SOLAR PV
0	156	0
1	156	0

2	156	0
3	156	0

Note that you need to put the column names in two sets of square brackets * the first set tells pandas you're passing a key / column name into the data frame, and * the second set of square brackets allows you to pass a *set* of column names.

1.4 Recap

- Pandas dataframes are sophisticated dicts of lists.
 - They have attributes like `columns` and `index` that have special meaning in the pandas context.
 - You can store any combination of data types in the dataframe
- You can access information in them as though they are dicts of lists