

Lecture 04.1 Pandas loc iloc

September 5, 2023

1 Pandas .loc and .iloc

Duncan Callaway

My objective in this notebook is to teach people how to access the information in a Pandas dataframe.

```
[1]: import pandas as pd
```

Let's load in the California ISO data we used last time

```
[2]: caiso_data = pd.read_csv('CAISO_2017to2018.csv')
```

```
[3]: caiso_data.head()
```

```
[3]:
```

	Unnamed: 0	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO	WIND	TOTAL	\
0	2017-08-29 00:00:00	1181	340	156	324		1551	
1	2017-08-29 01:00:00	1182	338	156	326		1556	
2	2017-08-29 02:00:00	1183	337	156	337		1325	
3	2017-08-29 03:00:00	1185	339	156	313		1158	
4	2017-08-29 04:00:00	1190	344	156	320		1209	

	SOLAR PV	SOLAR THERMAL
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

1.0.1 Q: What do you get if you call the dict of lists with a key?

```
[4]: caiso_data['BIOGAS']
```

```
[4]:
```

0	156
1	156
2	156
3	156
4	156
...	

```
8755    236
8756    234
8757    233
8758    234
8759    235
Name: BIOGAS, Length: 8760, dtype: int64
```

```
[5]: type(caiso_data['BIOGAS'])
```

```
[5]: pandas.core.series.Series
```

Ans: the list associated with the key. This is called a pandas ‘series’

1.0.2 Q: Figure out how to get solar production at 2pm on August 29 2017

First let’s check that we’ve got the right index for the time we want:

```
[6]: caiso_data['Unnamed: 0'][14]
```

```
[6]: '2017-08-29 14:00:00'
```

Now call the SOLAR PV column

```
[7]: caiso_data['SOLAR PV'][14]
```

```
[7]: 6820
```

1.0.3 Anatomy of the data frame.

Let’s talk a little about the anatomy of the data frame.

We have the following important attributes: 1. Rows 2. Columns 2. Index 3. Column names

The “index” can be numeric, but as we’ll see we can also make the indices strings.

```
[8]: caiso_data = pd.read_csv('CAISO_2017to2018.csv')
caiso_data.columns
```

```
[8]: Index(['Unnamed: 0', 'GEOTHERMAL', 'BIOMASS', 'BIOGAS', 'SMALL HYDRO',
          'WIND TOTAL', 'SOLAR PV', 'SOLAR THERMAL'],
          dtype='object')
```

Note that we can’t reassign easily because column and index names lists are immutable. Here is the workaround:

```
[9]: cols = caiso_data.columns.tolist()
cols[0] = 'Date and time'
caiso_data.columns = cols
caiso_data
```

```
[9]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS	SMALL HYDRO	\
0	2017-08-29 00:00:00	1181	340	156		324
1	2017-08-29 01:00:00	1182	338	156		326
2	2017-08-29 02:00:00	1183	337	156		337
3	2017-08-29 03:00:00	1185	339	156		313
4	2017-08-29 04:00:00	1190	344	156		320
...
8755	2018-08-28 19:00:00	962	332	236		581
8756	2018-08-28 20:00:00	967	336	234		547
8757	2018-08-28 21:00:00	972	336	233		502
8758	2018-08-28 22:00:00	975	333	234		361
8759	2018-08-28 23:00:00	977	333	235		262

	WIND TOTAL	SOLAR PV	SOLAR THERMAL
0	1551	0	0
1	1556	0	0
2	1325	0	0
3	1158	0	0
4	1209	0	0
...
8755	3300	70	24
8756	3468	0	17
8757	3310	0	17
8758	3068	0	0
8759	2921	0	0

```
[8760 rows x 8 columns]
```

Ok, that looks a little better for now.

As you can see, all the data are the same type of numeric value – MWh.

In these cases, sometimes it’s natural to “stack” the data.

We could do the stacking with a pandas command, `.stack`

1.1 Indexing and slicing in Pandas

First let’s figure out how to slice these data frames.

`.iloc` allows us to index and slice on integer row and column positions, like numpy:

```
[10]: caiso_data.iloc[1,1]
```

```
[10]: 1182
```

But what’s nice about `.iloc` is that you can also slice. It works just like numpy.

1.1.1 Q: Take a slice of the caiso_data dataframe that grabs the first four columns of data and the first 10 rows

```
[11]: caiso_data.iloc[0:10, :4]
```

```
[11]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS
0	2017-08-29 00:00:00	1181	340	156
1	2017-08-29 01:00:00	1182	338	156
2	2017-08-29 02:00:00	1183	337	156
3	2017-08-29 03:00:00	1185	339	156
4	2017-08-29 04:00:00	1190	344	156
5	2017-08-29 05:00:00	1194	351	157
6	2017-08-29 06:00:00	1196	359	155
7	2017-08-29 07:00:00	1194	363	153
8	2017-08-29 08:00:00	1187	364	153
9	2017-08-29 09:00:00	1189	367	157

1.1.2 Q: What would you do if you wanted to get the *last* 10 rows?

```
[12]: caiso_data.iloc[-10:, :4]
```

```
[12]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS
8750	2018-08-28 14:00:00	933	338	240
8751	2018-08-28 15:00:00	933	338	238
8752	2018-08-28 16:00:00	934	337	239
8753	2018-08-28 17:00:00	934	336	235
8754	2018-08-28 18:00:00	955	337	237
8755	2018-08-28 19:00:00	962	332	236
8756	2018-08-28 20:00:00	967	336	234
8757	2018-08-28 21:00:00	972	336	233
8758	2018-08-28 22:00:00	975	333	234
8759	2018-08-28 23:00:00	977	333	235

1.1.3 Q: Can you print out the last ten rows in reverse order?

```
[13]: caiso_data.iloc[: -10: -1, :4]
```

```
[13]:
```

	Date and time	GEO THERMAL	BIOMASS	BIOGAS
8759	2018-08-28 23:00:00	977	333	235
8758	2018-08-28 22:00:00	975	333	234
8757	2018-08-28 21:00:00	972	336	233
8756	2018-08-28 20:00:00	967	336	234
8755	2018-08-28 19:00:00	962	332	236
8754	2018-08-28 18:00:00	955	337	237
8753	2018-08-28 17:00:00	934	336	235
8752	2018-08-28 16:00:00	934	337	239
8751	2018-08-28 15:00:00	933	338	238

`.loc` is similar to `.iloc`, but it allows you to call the index and column names:

```
[14]: caiso_data.loc[0:5, 'Date and time']
```

```
[14]: 0    2017-08-29 00:00:00
      1    2017-08-29 01:00:00
      2    2017-08-29 02:00:00
      3    2017-08-29 03:00:00
      4    2017-08-29 04:00:00
      5    2017-08-29 05:00:00
      Name: Date and time, dtype: object
```

You can even slice with column names:

```
[15]: caiso_data.loc[0:5, 'Date and time': 'BIOGAS']
```

```
[15]:      Date and time  GEOTHERMAL  BIOMASS  BIOGAS
0  2017-08-29 00:00:00         1181      340     156
1  2017-08-29 01:00:00         1182      338     156
2  2017-08-29 02:00:00         1183      337     156
3  2017-08-29 03:00:00         1185      339     156
4  2017-08-29 04:00:00         1190      344     156
5  2017-08-29 05:00:00         1194      351     157
```

1.1.4 Q: Is `.loc` end-inclusive or exclusive when you slice?

Ans: *inclusive*. This is because it requires less knowledge about other rows in the DataFrame.

Note that this is true for both the index and the column names.

1.2 Recap

- Pandas dataframes are sophisticated dicts of lists.
 - They have attributes like `columns` and `index` that have special meaning in the pandas context.
 - You can store any combination of data types in the dataframe
- You can access information in them as though they are dicts of lists
- But you can also use the `.loc` and `.iloc` methods to access information in a way similar to numpy, including clean slicing.