

# Lecture 05.2 EDA

September 7, 2023

## 1 ER131: Data Cleaning and Exploratory Data Analysis

Duncan Callaway

In this notebook we'll work with PurpleAir data to explore the concepts of Structure, Granularity, Scope, Temporality and Faithfulness. Along the way we'll talk about data cleaning as well.

[Here's PurpleAir's website](#) – They have really cool maps!

The way I developed this lecture was by pulling the data down and exploring it. You'll see my (edited) process of examining the data.

This began by me visiting [this website](#) to look for data. I used the Chrome browser to pull data (other browsers didn't work).

The folks at PurpleAir also sent me a pdf describing their data, which is available from the instructors.

```
[1]: import numpy as np
import pandas as pd
import os
```

### 1.1 Structure: how are the data stored?

First let's look at what's in the data directory using `os.listdir` (remember this is a set of command line-style commands that work across platforms, i.e. mac, linux, windows)

```
[2]: os.listdir()
```

```
[2]: ['~$05 Pivot.pptx',
      '.DS_Store',
      'CAISO_2017to2018_stack.csv',
      'Lecture 05.1 Pivot.pdf',
      'Lecture 05.2 EDA.ipynb',
      '.ipynb_checkpoints',
      'US-EPA-PM2.5-AQI-Monitoring.png',
      'data',
      '05 Pivot.pptx',
      '05-06 Data Cleaning, EDA.pptx',
      '~$05-06 Data Cleaning, EDA.pptx',
      'Lecture 05.1 Pivot.ipynb']
```

```
[3]: !ls # this does the same thing
```

```
05 Pivot.pptx          Lecture 05.2 EDA.ipynb
05-06 Data Cleaning, EDA.pptx  US-EPA-PM2.5-AQI-Monitoring.png
CAISO_2017to2018_stack.csv    data
Lecture 05.1 Pivot.ipynb     ~$05 Pivot.pptx
Lecture 05.1 Pivot.pdf       ~$05-06 Data Cleaning, EDA.pptx
```

Let's look in the data directory:

```
[4]: os.listdir('data')
```

```
[4]: ['Alameda Gold Coast (outside) (37.767347 -122.267255) Primary Real Time
09_08_2021 09_07_2022.csv',
      'Bower House (outside) (37.803884 -122.297151) Primary Real Time 09_08_2021
09_07_2022.csv',
      'Backyard (outside) (37.826875 -122.245254) Primary Real Time 09_08_2021
09_07_2022.csv',
      'Moraga Ave (outside) (37.83023 -122.239963) Primary Real Time 09_08_2021
09_07_2022.csv',
      'manzanita at villanova (outside) (37.84099 -122.196456) Primary Real Time
09_08_2021 09_07_2022.csv',
      'B59-Mech (outside) (37.875921 -122.253082) Primary Real Time 09_08_2021
09_07_2022.csv']
```

### 1.1.1 Q: What can we learn from these file names?

- the sensor location appears to be provided in lat / lon coordinates in parens
- the date range is listed
- they are probably csv files.

If you type the lat-lon values into google maps, you'll find they correspond to the locations of purple air sensors with the same name. [Here](#) is a route through these sites.

Before proceeding let's find the size of some of these files:

```
[6]: os.path.getsize('data/Bower House (outside) (37.803884 -122.297151) Primary_
↳Real Time 09_08_2021 09_07_2022.csv')
```

```
[6]: 28354380
```

What are the units? Let's shift tab in to `getsize` to find out.

```
[9]: os.path.getsize
```

```
[9]: <function genericpath.getsize(filename)>
```

Not much information. Google search reveals [this](#) information page, which says the units are bytes.

```
[10]: 1e-6*os.path.getsize('data/Bower House (outside) (37.803884 -122.297151)_
↳Primary Real Time 09_08_2021 09_07_2022.csv')
```

```
[10]: 28.35438
```

SO 28 Mb.

Let's read in one of the .csv files:

```
[11]: Bower = pd.read_csv('data/Bower House (outside) (37.803884 -122.297151) Primary_
↳Real Time 09_08_2021 09_07_2022.csv')
```

```
[12]: Bower.head()
```

```
[12]:
```

	created_at	entry_id	PM1.0_CF1_ug/m3	PM2.5_CF1_ug/m3	\
0	2021-09-08 00:00:30 UTC	271065	14.34	29.12	
1	2021-09-08 00:02:30 UTC	271066	15.74	29.59	
2	2021-09-08 00:04:30 UTC	271067	13.86	28.89	
3	2021-09-08 00:06:30 UTC	271068	14.43	28.59	
4	2021-09-08 00:08:30 UTC	271069	13.00	27.02	

  

	PM10.0_CF1_ug/m3	UptimeMinutes	RSSI_dbm	Temperature_F	Humidity_%	\
0	38.14	40112.0	-71.0	84.0	38.0	
1	37.28	40114.0	-70.0	84.0	38.0	
2	41.70	40116.0	-75.0	84.0	37.0	
3	45.93	40118.0	-76.0	84.0	37.0	
4	41.47	40120.0	-70.0	84.0	36.0	

  

	PM2.5_ATM_ug/m3	Unnamed: 10
0	28.73	NaN
1	29.15	NaN
2	28.64	NaN
3	28.41	NaN
4	27.02	NaN

### 1.1.2 Q: What do you notice about the file contents?

Several things to ask from this: 1. Dates are UTC. 2. Each entry has a unique ID – could be used to check for time stamp errors or gaps in data 3. Headers have ‘CF1’ or ‘ATM’ at the top – what does that mean? 1. From the PurpleAir documentation, in this directory, “ATM is”atmospheric”, meant to be used for outdoor applications. CF=1 is meant to be used for indoor or controlled environment applications. However, PurpleAir uses CF=1 values on the map. This value is lower than the ATM value in higher measured concentrations.”

2. The explanation is a little vague and suggests further exploration required! 3. This cool paper suggests that the ATM data are ‘raw’ measurements and that CF\_1 data have a 3/2 multiplication at concentrations over  $25 \mu\text{g} / \text{m}^3$  4. The columns “UptimeMinutes” and “RSSI\_dbm” are not immediately obvious 1. again from documentation: “uptimeminutes” is time since last restart, and “RSSI\_dbm” is wifi signal strength for the device.

5. The “unnamed: 10” column seems useless, why is it there? 1. Looking at the data we see a comma before the \n (newline character) at the end of the first (header) line, it appears this is generating the extra row.

```
[14]: N = 2
airdat = "data/Bower House (outside) (37.803884 -122.297151) Primary Real Time_
↪09_08_2021 09_07_2022.csv"
with open(airdat) as myfile:
    head = [next(myfile) for x in range(N)]
print(head)
```

```
['created_at,entry_id,PM1.0_CF1_ug/m3,PM2.5_CF1_ug/m3,PM10.0_CF1_ug/m3,UptimeMin
utes,RSSI_dbm,Temperature_F,Humidity_%,PM2.5_ATM_ug/m3,\n', '2021-09-08 00:00:30
UTC,271065,14.34,29.12,38.14,40112.00,-71.00,84.00,38.00,28.73\n']
```

```
[15]: Bower.describe()
```

```
[15]:
```

	entry_id	PM1.0_CF1_ug/m3	PM2.5_CF1_ug/m3	PM10.0_CF1_ug/m3	\
count	351942.000000	351942.000000	351942.000000	351942.000000	
mean	390592.636500	7.703851	15.989965	24.196788	
std	68985.096555	9.894243	19.466634	24.491579	
min	271065.000000	0.000000	0.000000	0.000000	
25%	336116.000000	1.210000	3.590000	6.750000	
50%	380108.500000	3.760000	8.750000	16.410000	
75%	444398.750000	10.220000	20.280000	33.270000	
max	532384.000000	336.360000	1297.910000	1582.450000	

	UptimeMinutes	RSSI_dbm	Temperature_F	Humidity_%	\
count	351942.000000	351942.000000	351942.000000	351942.000000	
mean	26747.629834	-70.180916	65.679018	50.249541	
std	18879.817383	5.026611	11.876395	14.772064	
min	1.000000	-93.000000	36.000000	5.000000	
25%	10005.000000	-71.000000	58.000000	41.000000	
50%	23995.500000	-69.000000	63.000000	55.000000	
75%	41299.000000	-67.000000	71.000000	62.000000	
max	71582.000000	-53.000000	126.000000	77.000000	

	PM2.5_ATM_ug/m3	Unnamed: 10
count	351942.000000	0.0
mean	14.203361	NaN
std	14.610303	NaN
min	0.000000	NaN
25%	3.590000	NaN
50%	8.750000	NaN
75%	20.270000	NaN
max	865.120000	NaN

As you can see, at this location the average CF1 value is more than the EPA standards. As an aside, before we do more EDA, let's check the other location.

```
[17]: backyard = pd.read_csv('data/Backyard (outside) (37.826875 -122.245254) Primary_
↪Real Time 09_08_2021 09_07_2022.csv')
```

```
np.mean(backyard['PM2.5_CF1_ug/m3'])
```

```
[17]: 11.776786947061435
```

```
[18]: moraga = pd.read_csv('data/Moraga Ave (outside) (37.83023 -122.239963) Primary_␣  
↪Real Time 09_08_2021 09_07_2022.csv')  
np.mean(moraga['PM2.5_CF1_ug/m3'])
```

```
[18]: 10.18030246986375
```

```
[19]: manzanita = pd.read_csv('data/manzanita at villanova (outside) (37.84099 -122.  
↪196456) Primary Real Time 09_08_2021 09_07_2022.csv')  
np.mean(manzanita['PM2.5_CF1_ug/m3'])
```

```
[19]: 8.031333700168869
```

```
[20]: alameda = pd.read_csv('data/Alameda Gold Coast (outside) (37.767347 -122.  
↪267255) Primary Real Time 09_08_2021 09_07_2022.csv')  
np.mean(alameda['PM2.5_CF1_ug/m3'])
```

```
[20]: 13.041692402454741
```

Now you can see that the mean PM2.5 numbers vary significantly by location.

If you inspect the data, you'll see a general trend: the further away from the Bay the sensor is, the lower its mean.

Let's dig in to one sensor a little more

## 1.2 Granularity: how are the data aggregated?

We'll talk a little more about Temporality in a moment, but time also matters for thinking about granularity.

First we need to pay attention to the fact that this is UTC. Let's put it in datetime format to prevent mistakes.

```
[21]: Bowertime = pd.to_datetime(Bower['created_at'], utc=True)
```

```
[22]: Bower['created_at']=Bowertime
```

```
[23]: Bower['created_at'].dtype
```

```
[23]: datetime64[ns, UTC]
```

Yes, that response really means the time are recorded down to the nanosecond.

Note: The data are instantaneous measurements, not averaged over time.

\* So these data have granularity of nanoseconds! \* In practice, this just means there is *no* aggregation in the primary data.

```
[24]: Bower.head()
```

```
[24]:
```

	created_at	entry_id	PM1.0_CF1_ug/m3	PM2.5_CF1_ug/m3	\
0	2021-09-08 00:00:30+00:00	271065	14.34	29.12	
1	2021-09-08 00:02:30+00:00	271066	15.74	29.59	
2	2021-09-08 00:04:30+00:00	271067	13.86	28.89	
3	2021-09-08 00:06:30+00:00	271068	14.43	28.59	
4	2021-09-08 00:08:30+00:00	271069	13.00	27.02	

  

	PM10.0_CF1_ug/m3	UptimeMinutes	RSSI_dbm	Temperature_F	Humidity_%	\
0	38.14	40112.0	-71.0	84.0	38.0	
1	37.28	40114.0	-70.0	84.0	38.0	
2	41.70	40116.0	-75.0	84.0	37.0	
3	45.93	40118.0	-76.0	84.0	37.0	
4	41.47	40120.0	-70.0	84.0	36.0	

  

	PM2.5_ATM_ug/m3	Unnamed: 10
0	28.73	NaN
1	29.15	NaN
2	28.64	NaN
3	28.41	NaN
4	27.02	NaN

Nice thing about the datetime formate is that you can easily get time information out of it. For example let's look at the 1,000th entry:

```
[38]: Bower.iloc[1000,0].hour
```

```
[38]: 9
```

Note, we could rename the cols to make things easier if we wished. I'm not going to because we're not going to be workign with this data set for long, but in other cases you might decide to.

### 1.3 Scope: how much time, how many people, what spatial area?

So far we have focused on data from one location – A sensor in West Oakland.

From the file name it looks like the time is from the last 12 months, let's confirm:

```
[39]: Bower['created_at'].min()
```

```
[39]: Timestamp('2021-09-08 00:00:30+0000', tz='UTC')
```

```
[40]: Bower['created_at'].max()
```

```
[40]: Timestamp('2022-09-07 23:58:17+0000', tz='UTC')
```

So it's about one year of data.

Does the data cover the topic of interest?

In this case, we need to answer the question: For the PurpleAir data, what topic of interest might the data cover?

→ **class discussion on this.** Possible answers why the data might be of interest \* near highways and port of oakland \* near communities that are historically underserved

Possible reasons *not* of interest: \* more important to look at many recent wildfire seasons \* it might be valuable to compare across sites rather than evaluate just one.

## 1.4 Temporality: How is time represented in the data?

We've already figured out that we're working with UTC dates. UTC is "universal time coordinated" and is essentially greenwich mean time, the time on the prime meridian.

Can we figure out how frequent measurements are?

Unfortunately I found it difficult to take differences with datetime objects, so I had to write a for loop:

```
[77]: diffs = np.zeros(len(Bower['created_at']))

for i in range(0, len(diffs)-1):
    diffs[i] = ((Bower['created_at'][i+1]
                 - Bower['created_at'][i]).total_seconds()) # we apply
    ↪total_seconds in order to store the data as a float in the list

diffs = np.sort((diffs))

print('max diffs:', diffs[:-30:-1])
print('median:', np.median(diffs))
```

```
max diffs: [79585. 43629. 3842. 1782. 1782. 1561. 1443. 1441. 1437.
976.
          976.  727.  724.  720.  718.  718.  495.  495.  480.  480.
          388.  363.  360.  359.  358.  358.  346.  292.  276.]
median: 120.0
```

Looks like for the most part we're sampling every 2 minutes, with a few gaps in the data.

## 1.5 Faithfulness: are the data trustworthy?

This one's much harder to assess. Let's have a look at some basic things we might care about

```
[78]: sum(Bower['PM2.5_ATM_ug/m3'].isna())
```

```
[78]: 0
```

That tells us there are no NaN values in the PM2.5 data. Impressive!

```
[79]: Bower.describe()
```

```
[79]:
```

	entry_id	PM1.0_CF1_ug/m3	PM2.5_CF1_ug/m3	PM10.0_CF1_ug/m3	\
count	351942.000000	351942.000000	351942.000000	351942.000000	
mean	390592.636500	7.703851	15.989965	24.196788	
std	68985.096555	9.894243	19.466634	24.491579	
min	271065.000000	0.000000	0.000000	0.000000	
25%	336116.000000	1.210000	3.590000	6.750000	
50%	380108.500000	3.760000	8.750000	16.410000	
75%	444398.750000	10.220000	20.280000	33.270000	
max	532384.000000	336.360000	1297.910000	1582.450000	

  

	UptimeMinutes	RSSI_dbm	Temperature_F	Humidity_%	\
count	351942.000000	351942.000000	351942.000000	351942.000000	
mean	26747.629834	-70.180916	65.679018	50.249541	
std	18879.817383	5.026611	11.876395	14.772064	
min	1.000000	-93.000000	36.000000	5.000000	
25%	10005.000000	-71.000000	58.000000	41.000000	
50%	23995.500000	-69.000000	63.000000	55.000000	
75%	41299.000000	-67.000000	71.000000	62.000000	
max	71582.000000	-53.000000	126.000000	77.000000	

  

	PM2.5_ATM_ug/m3	Unnamed: 10
count	351942.000000	0.0
mean	14.203361	NaN
std	14.610303	NaN
min	0.000000	NaN
25%	3.590000	NaN
50%	8.750000	NaN
75%	20.270000	NaN
max	865.120000	NaN

That's a pretty high PM2.5 average. And the max is very suspiciously high. What's going on?

Options: 1. Wildfire smoke really pumped up the 2.5 values 2. We have a lot of missing data and only values during the wild fires 3. There are some erroneously high values.

Let's start by looking at how many values are big.

```
[81]: log_ind = Bower.loc[:, 'PM2.5_CF1_ug/m3'] > 500 # gives a list for logical_
      ↪ indexing
      Bower.loc[log_ind, 'PM2.5_CF1_ug/m3']
```

```
[81]: 150102      608.21
      150822      608.21
      305431     1297.91
      305432      878.52
      Name: PM2.5_CF1_ug/m3, dtype: float64
```

Let's look in the vicinity of the high values to see if we believe the trend:



[82]: Bower.loc[305420:305435,:]

```
[82]:
```

		created_at	entry_id	PM1.0_CF1_ug/m3	PM2.5_CF1_ug/m3	\
305420	2022-07-05	05:46:09+00:00	485863	5.17	10.33	
305421	2022-07-05	05:48:09+00:00	485864	6.66	10.88	
305422	2022-07-05	05:50:09+00:00	485865	6.81	12.11	
305423	2022-07-05	05:52:09+00:00	485866	6.98	11.89	
305424	2022-07-05	05:54:09+00:00	485867	5.82	11.07	
305425	2022-07-05	05:56:09+00:00	485868	5.65	10.33	
305426	2022-07-05	05:58:09+00:00	485869	7.17	13.38	
305427	2022-07-05	06:00:09+00:00	485870	5.09	8.67	
305428	2022-07-05	06:02:09+00:00	485871	4.52	8.33	
305429	2022-07-05	06:04:09+00:00	485872	5.62	10.02	
305430	2022-07-05	06:06:09+00:00	485873	11.17	17.88	
305431	2022-07-05	06:08:09+00:00	485874	336.36	1297.91	
305432	2022-07-05	06:10:09+00:00	485875	332.55	878.52	
305433	2022-07-05	06:12:09+00:00	485876	173.87	342.52	
305434	2022-07-05	06:14:09+00:00	485877	68.84	129.05	
305435	2022-07-05	06:16:09+00:00	485878	26.69	48.02	

	PM10.0_CF1_ug/m3	UptimeMinutes	RSSI_dbm	Temperature_F	Humidity_%	\
305420	12.72	70909.0	-67.0	73.0	59.0	
305421	12.31	70911.0	-67.0	72.0	59.0	
305422	14.12	70913.0	-72.0	72.0	60.0	
305423	14.58	70915.0	-69.0	72.0	60.0	
305424	14.45	70917.0	-68.0	73.0	60.0	
305425	13.14	70919.0	-67.0	73.0	60.0	
305426	16.48	70921.0	-72.0	72.0	60.0	
305427	10.78	70923.0	-65.0	72.0	60.0	
305428	11.19	70925.0	-69.0	72.0	59.0	
305429	11.29	70927.0	-67.0	72.0	60.0	
305430	23.90	70929.0	-71.0	72.0	60.0	
305431	1582.45	70931.0	-68.0	72.0	60.0	
305432	1027.02	70933.0	-71.0	72.0	60.0	
305433	392.19	70935.0	-68.0	68.0	59.0	
305434	155.05	70937.0	-67.0	73.0	60.0	
305435	60.86	70939.0	-69.0	72.0	60.0	

	PM2.5_ATM_ug/m3	Unnamed: 10
305420	10.33	NaN
305421	10.88	NaN
305422	12.11	NaN
305423	11.89	NaN
305424	11.07	NaN
305425	10.33	NaN
305426	13.38	NaN
305427	8.67	NaN

305428	8.33	NaN
305429	10.02	NaN
305430	17.88	NaN
305431	865.12	NaN
305432	585.00	NaN
305433	227.58	NaN
305434	85.29	NaN
305435	39.97	NaN

Looks like there was a stretch of time with really high values, somewhat suspiciously clustered around 5000. If I were doing more work here I would look into the sensor more carefully to see if there is any significance to that number.

But for now – let’s just go ahead and drop them and see what happens:

```
[83]: Bower.loc[log_ind, 'PM2.5_CF1_ug/m3'] = np.nan
      Bower.describe()
```

```
[83]:
```

	entry_id	PM1.0_CF1_ug/m3	PM2.5_CF1_ug/m3	PM10.0_CF1_ug/m3	\
count	351942.000000	351942.000000	351938.000000	351942.000000	
mean	390592.636500	7.703851	15.980506	24.196788	
std	68985.096555	9.894243	19.240002	24.491579	
min	271065.000000	0.000000	0.000000	0.000000	
25%	336116.000000	1.210000	3.590000	6.750000	
50%	380108.500000	3.760000	8.750000	16.410000	
75%	444398.750000	10.220000	20.280000	33.270000	
max	532384.000000	336.360000	346.050000	1582.450000	

	UptimeMinutes	RSSI_dbm	Temperature_F	Humidity_%	\
count	351942.000000	351942.000000	351942.000000	351942.000000	
mean	26747.629834	-70.180916	65.679018	50.249541	
std	18879.817383	5.026611	11.876395	14.772064	
min	1.000000	-93.000000	36.000000	5.000000	
25%	10005.000000	-71.000000	58.000000	41.000000	
50%	23995.500000	-69.000000	63.000000	55.000000	
75%	41299.000000	-67.000000	71.000000	62.000000	
max	71582.000000	-53.000000	126.000000	77.000000	

	PM2.5_ATM_ug/m3	Unnamed: 10
count	351942.000000	0.0
mean	14.203361	NaN
std	14.610303	NaN
min	0.000000	NaN
25%	3.590000	NaN
50%	8.750000	NaN
75%	20.270000	NaN
max	865.120000	NaN

You can see the average came down a little, and the standard deviation came *really* far down. And

as we'd hope the max is now below 500.

[ ]: